

Assignment 2: An Investigation of the Components of Pre-Training and Post-Training in LLMs

Course: Foundation for Large Language Models (FLLM - AI60213)

Due Date: October-30-2025

[Total marks: 20]

Advisory

Using GenAI tools for solving assignments is discouraged, as this will increase the probability of plagiarism with others who may choose the same path.

Submission Instructions

1. Only one member from each group needs to submit the assignment through Moodle to avoid any duplicate submissions.
2. Submit the solution file as: `Group(number)_assignment2_solution.ipynb` and also include a corresponding `.py` file.

Plagiarism Policy

1. Code plagiarism will be checked on the designated blocks which the students have to code.
2. No penalty till 30% (Baseline).
3. Each 10% increase over baseline will cost 1 mark. Beyond 70%, absolute -2 (negative).
4. Two groups having similar code will be penalized with the same amount. No arbitration will be done.
5. Late submission policy: One week window with a penalty of 3 marks.

General Instructions

This assignment is designed to give you a hands-on understanding of the components of Pre-Training and Post-Training in LLMs. You are required to do experiments using **BART-base** [3] model and two types of adapters - residual and sparse for the multi-document summarization (MDS) [1] task.

You should submit a single PDF report containing your answers to the questions, including any plots and tables. Your report should clearly explain your methodology, observations, and conclusions for each question. You should also submit your source code (e.g., a Jupyter Notebook file) with clear comments.

For your experiments, you can use the Hugging Face ‘transformers’ and ‘datasets’ libraries. The ‘facebook/bart-base’ model and ‘allenai/mslr2022’ dataset from Huggingface should be sufficient for all parts and can be run in a Google Colab environment.

(Note: Please note that all the hyperparameters are already defined in the provided code file. Kindly use those parameters while executing the assignment. If you wish to experiment with different hyperparameters for higher performance, you may do so — but make sure to clearly mention the changes and provide justification for each modification in your submission report.)

Question 1: Analyzing with or without Adapter fine-tuning for Multi-Document Summarization (MDS) task.

[Marks: 20 (10 + 2 + 4 + 4)]

Background:

Multi-document summarization (MDS) aims to generate a single, coherent summary from multiple related documents that discuss the same topic. The MS² [1] dataset (Multi-Document Summarization of Medical Studies - ‘allenai/mslr2022’) is designed specifically for this purpose. It consists of groups of biomedical study abstracts, review background, and target (summary). In this question, you are required to demonstrate your understanding of adapter-based fine-tuning and full fine-tuning on a small subset of the MS² dataset (the first 2000 samples from the training set for all settings and the first 500 samples from the validation set).

Tasks:

1. Choose either the **residual** [4] or the **sparse** [2] adapter. The BOTTLEBECK_MAP sets 128 for the **residual** adapter and 256 for the **sparse** adapter. For the **summarization** task, the MAXLEN_MAP is 1024 tokens for the source and 128 tokens for the target. Use the **facebook/bart-base** model as a base model for adapter-based fine-tuning, following the first 2000 samples of the training set. Save the fine-tuned model after training. Follow the hyperparameters given in Table 1.
2. Using the saved fine-tuned model, perform multi-document summarization on the first 500 samples of the validation set.

3. Using the generated summaries and the given target summaries of the first 500 samples of the validation set, compute the **rouge_score**. Report the average scores for ROUGE-1, ROUGE-2, and ROUGE-L, including precision, recall, and F1-score.
4. Perform full fine-tuning of the base model (without using adapters) with the same hyper-parameters given in Table 1. Compute the same ROUGE metrics as in part 3 and compare the results with the adapter-based model results. Discuss your findings and highlight any differences observed between the two settings.

Table 1: Hyperparameters and their values

Hyperparameter	Value
Optimizer	AdamW
Learning Rate	$3e-3$ or $5e-5$
Epochs	5 or 8
Batch Size	8 or 16 or 32
BOTTLEBECK_MAP Residual Adapter	128
BOTTLEBECK_MAP Sparse Adapter	256

References

- [1] Jay DeYoung, Iz Beltagy, Madeleine van Zuylen, Bailey Kuehl, and Lucy Lu Wang. MS²: Multi-document summarization of medical studies. In Marie-Francine Moens, Xuanjing Huang, Lucia Specia, and Scott Wen-tau Yih, editors, *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 7494–7513, Online and Punta Cana, Dominican Republic, November 2021. Association for Computational Linguistics. doi: 10.18653/v1/2021.emnlp-main.594. URL <https://aclanthology.org/2021.emnlp-main.594/>.
- [2] Shwai He, Liang Ding, Daize Dong, Miao Zhang, and Dacheng Tao. Sparseadapter: An easy approach for improving the parameter-efficiency of adapters. *arXiv preprint arXiv:2210.04284*, 2022.
- [3] Mike Lewis, Yinhan Liu, Naman Goyal, Marjan Ghazvininejad, Abdelrahman Mohamed, Omer Levy, Ves Stoyanov, and Luke Zettlemoyer. Bart: Denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension. *arXiv preprint arXiv:1910.13461*, 2019.
- [4] Zhaojiang Lin, Andrea Madotto, and Pascale Fung. Exploring versatile generative language model via parameter-efficient transfer learning. *arXiv preprint arXiv:2004.03829*, 2020.